



UNIVERSITY OF PADOVA

Department of Physics and Astronomy 'Galileo Galilei'

MSc in Physics of Data

Streaming processing of cosmic rays using Drift Tubes detectors

Gabriele Bertinelli - Roben Bhatti - Diego Bonato - Martina Cacciola

July 12th, 2023

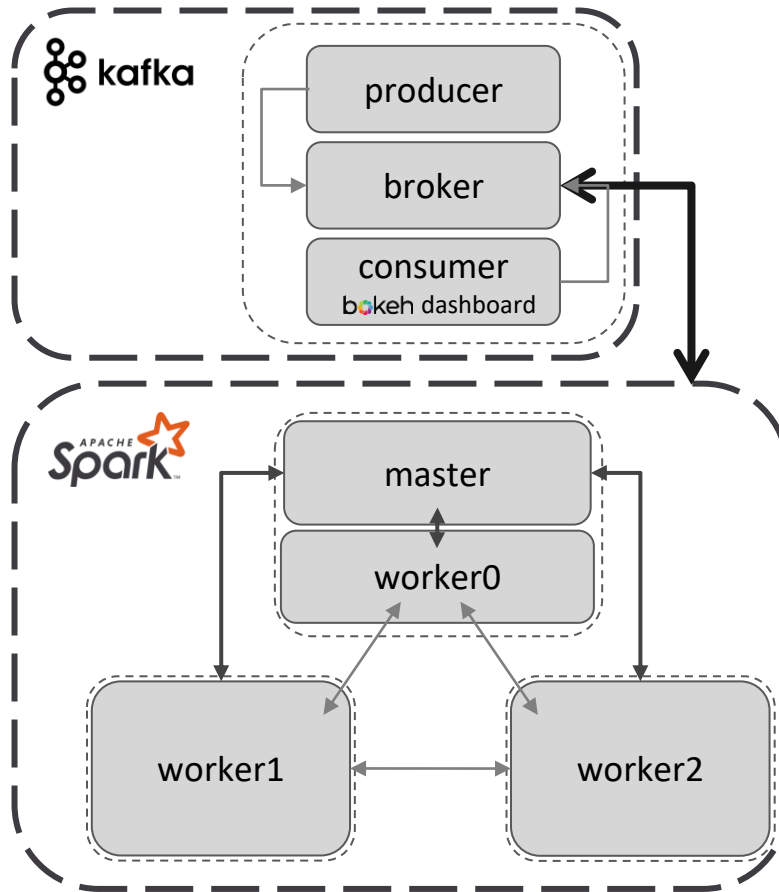
Goal of the project

- Our project aims to create a simulated data processing network for a particle physics detector, enabling **real-time** monitoring of the results through a dashboard.
- The dataset consists of multiple text files in comma-separated values (CSV) format, hosted on Cloud Veneto's cloud storage bucket.
- Our objective is to simulate a continuous DAQ stream by injecting the provided dataset into a Kafka **topic**. After performing data cleansing, we package the extracted information into individual **messages** per batch and inject them into a new Kafka topic.

Tools

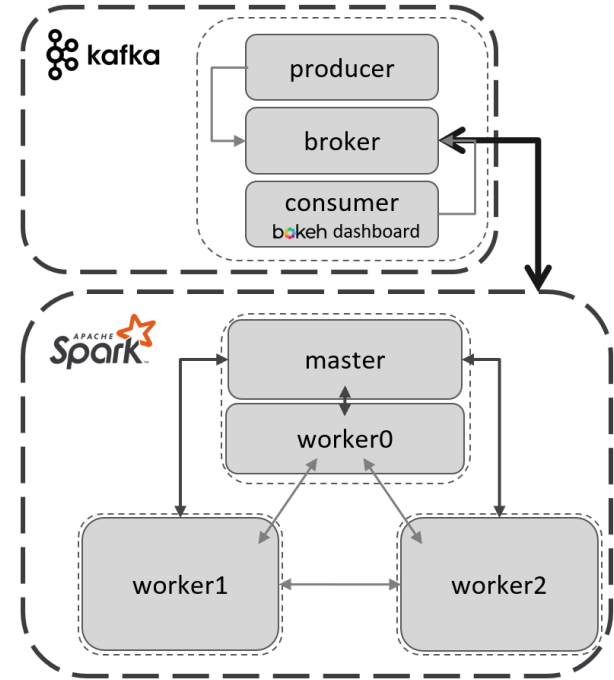
To implement our project, we utilize various frameworks and Python packages, including:

- **Kafka 3.3.2:** This distributed event streaming platform is employed to manage the data streaming process. We interface with Kafka using the *kafka* package.
- **Spark 3.3.2:** This cluster computing framework is used for data analytics, enabling distributed computation.
- **Bokeh 3.1.1:** This Python package is utilized to enhance the visual representation of our data through interactive plotting.



Network architecture

- Spark master → 2 cores
- worker0 → 2 cores } 8 GB RAM
- worker1 → } 4 cores
- worker2 → } 8 GB RAM
- Kafka broker } 4 cores
- producer } 8 GB RAM
- consumer }



Producer

- Connect to the cluster → `KAFKA_BOOTSTRAP_SERVER`
- Define a producer
- Create two topics → `data_raw`, `data_clean`
 - `num_partitions`
 - `replicaton_factor (=1)`

```
data_raw = NewTopic(name='data_raw',  
                    num_partitions=100,  
                    replication_factor=1)  
  
data_clean = NewTopic(name='data_clean',  
                      num_partitions=1,  
                      replication_factor=1)  
  
kafka_admin.create_topics(new_topics=[data_raw, data_clean])
```

Producer

- Connect to s3 bucket
- Access each file by ``Key`` value → save it in a Pandas Dataframe
- Loop in the dataframe → append each row to a message
- Send message asynchronously → `KafkaProducer.send()`
- Full batch → `KafkaProducer.flush()`

```
df=pd.read_csv(s3_client.get_object(Bucket='mapd-minidt-stream', Key=obj['Key']).get('Body'), sep=' ')
'...'
    # append a record to the msg to send
    producer.send('data_raw', row)
'...'

if batch_counter==batch_size:
    producer.flush()
```

Spark analysis

- Session and Context creation
- To test performance → we varied these parameters
 - `spark.executor.instances`: n° of executors
 - `spark.executor.cores`: n° of CPU cores for each executor
 - `spark.sql.shuffle.partitions`: n° partitions used when shuffling for joins/aggregations
 - `spark.sql.execution.arrow.pyspark.enabled`: in memory columnar format →
no major differences → left `true``
- Producer creation → send the final message to the `data_clean` topic

Spark analysis

```
# read streaming df from kafka
```

```
inputDF = spark\  
  .readStream\  
  .format("kafka")\  
  .option("kafka.bootstrap.servers", KAFKA_BOOTSTRAP_SERVER)\  
  .option("kafka.consumer.pollTimeoutMs", 1000)\  
  .option('subscribe', 'data_raw')\  
  .option("startingOffsets", "latest") \  
  .load()
```

```
# extract the value from the kafka message
```

```
rawraw_data = inputDF.select(col("value").cast("string")).alias("csv")
```

```
# split the csv line in the corresponding fields
```

```
raw_data = rawraw_data.selectExpr("cast(split(value, ',')[0] as int) as HEAD",  
  "cast(split(value, ',')[1] as int) as FPGA",  
  "cast(split(value, ',')[2] as int) as TDC_CHANNEL",  
  "cast(split(value, ',')[3] as long) as ORBIT_CNT",  
  "cast(split(value, ',')[4] as int) as BX_COUNTER",  
  "cast(split(value, ',')[5] as double) as TDC_MEAS")
```

- Select useful data, add `CHAMBER` column

Spark analysis

- Create a function to be applied to each batch, `batch_processing`
- Tasks:
 - total number of processed hits: ``hit_count``
 - total number of processed hits per chamber: ``hit_count_chamber``
 - histogram of the counts of active ``TDC_CHANNEL``, per chamber: ``ch*_tdc_counts_list``
 - histogram of the total number of active ``TDC_CHANNEL``
in each ``ORBIT_CNT``, per chamber: ``ch*_orbit_counts_list``
- Create JSON message → send to topic

Spark analysis

```
def batch_processing(df, epoch_id):  
  
    df = df.persist()  
  
    # 1: total number of processed hits, post-cleansing (1 value per batch)  
    hit_count = df.count()  
  
    # 2: total number of processed hits, post-cleansing, per chamber (4 values per batch)  
    hit_count_chamber = df.groupby('CHAMBER').agg(count('TDC_CHANNEL').alias('HIT_COUNT'))\  
        .sort("CHAMBER").select('HIT_COUNT')\  
        .agg(collect_list('HIT_COUNT')).collect()
```

Spark analysis

```
def batch_processing(df, epoch_id):  
    '...'  
  
    # 3: histogram of the counts of active TDC_CHANNEL, per chamber (4 arrays per batch)  
    tdc_counts = df.groupby(['CHAMBER', 'TDC_CHANNEL']).agg(count('TDC_CHANNEL').alias('TDC_COUNTS'))  
    tdc_counts = tdc_counts.persist()  
    # Filter the tdc_counts DataFrame for each chamber  
    ch*_tdc_counts = tdc_counts.filter(tdc_counts.CHAMBER == *).select('TDC_CHANNEL', 'TDC_COUNTS')\  
        .sort("TDC_CHANNEL").toPandas()  
  
    #Save it in a list  
    ch*_tdc_channels_list = list(ch*_tdc_counts['TDC_CHANNEL'])  
    ch*_tdc_counts_list   = list(ch*_tdc_counts['TDC_COUNTS'])  
  
    # 4: histogram of the total number of active TDC_CHANNEL in each ORBIT_CNT, per chamber (4 arrays per batch)  
    orbit_count=df.groupby(['CHAMBER', 'ORBIT_CNT']).agg(countDistinct("TDC_CHANNEL").alias('TDC_ORBIT'))  
    orbit_count = orbit_count.persist()  
    ch*_orbit_counts = orbit_count.filter(orbit_count.CHAMBER == *).select('ORBIT_CNT', 'TDC_ORBIT')\  
        .sort("ORBIT_CNT").toPandas()  
  
    #Save it in a list  
    ch*_orbit_list      = list(ch*_orbit_counts['ORBIT_CNT'])  
    ch*_orbit_counts_list = list(ch*_orbit_counts['TDC_ORBIT'])  
  
    df.unpersist()  
    tdc_counts.unpersist()  
    orbit_count.unpersist()
```

Spark analysis

```
def batch_processing(df, epoch_id):  
    ...  
    msg = { 'msg_ID': ID,  
            'hit_count': hit_count,  
            'hit_count_chamber': hit_count_chamber[0][0],  
  
            'tdc_counts_chamber': {  
                '0': {  
                    'bin_edges': ch0_tdc_channels_list,  
                    'hist_counts': ch0_tdc_counts_list  
                },  
                ...  
            },  
            'active_tdc_chamber': {  
                '0': {  
                    'bin_edges': ch0_orbit_list,  
                    'hist_counts': ch0_orbit_counts_list  
                },  
                ...  
            }  
        }  
    producer.send('data_clean', json.dumps(msg).encode('utf-8'))  
    producer.flush()
```

Spark analysis

- Apply the function to the streaming dataset

```
query = raw_data.writeStream\  
    .outputMode("update")\  
    .foreachBatch(batch_processing)\  
    .option("checkpointLocation", "checkpoint")\  
    .trigger(processingTime='5 seconds')\  
    .start()  
query.awaitTermination(5)
```

Consumer - Dashboard

- Creation of the consumer

```
# Initialize Kafka consumer by subscribing to the topic
consumer = KafkaConsumer('data_clean',
                          bootstrap_servers = KAFKA_BOOTSTRAP_SERVER)
```

- Poll messages

```
def polling():
    for msg in consumer:
        value = json.loads(msg.value.decode('utf-8')) # Convert the message value to a dictionary
        break
    return value

def create_value():
    #instantiate dictionary
    combined_dict = {}
    poll = polling()

    x = poll["msg_ID"]
    y = poll['hit_count']
    combined_dict["p1"] = {"x": [x], "y": [y]}
    ...
    return combined_dict
```

Consumer - Dashboard

- Create periodic function

```
def update():  
    #poll data and filter them  
    new_data = create_value()  
  
    #stream new data to dashboard  
    source1.stream(new_data["p1"], rollover=10)  
    source2.stream(new_data["p2_0"], rollover=10)  
  
    '...'
```

- Layout configuration and axis formatting

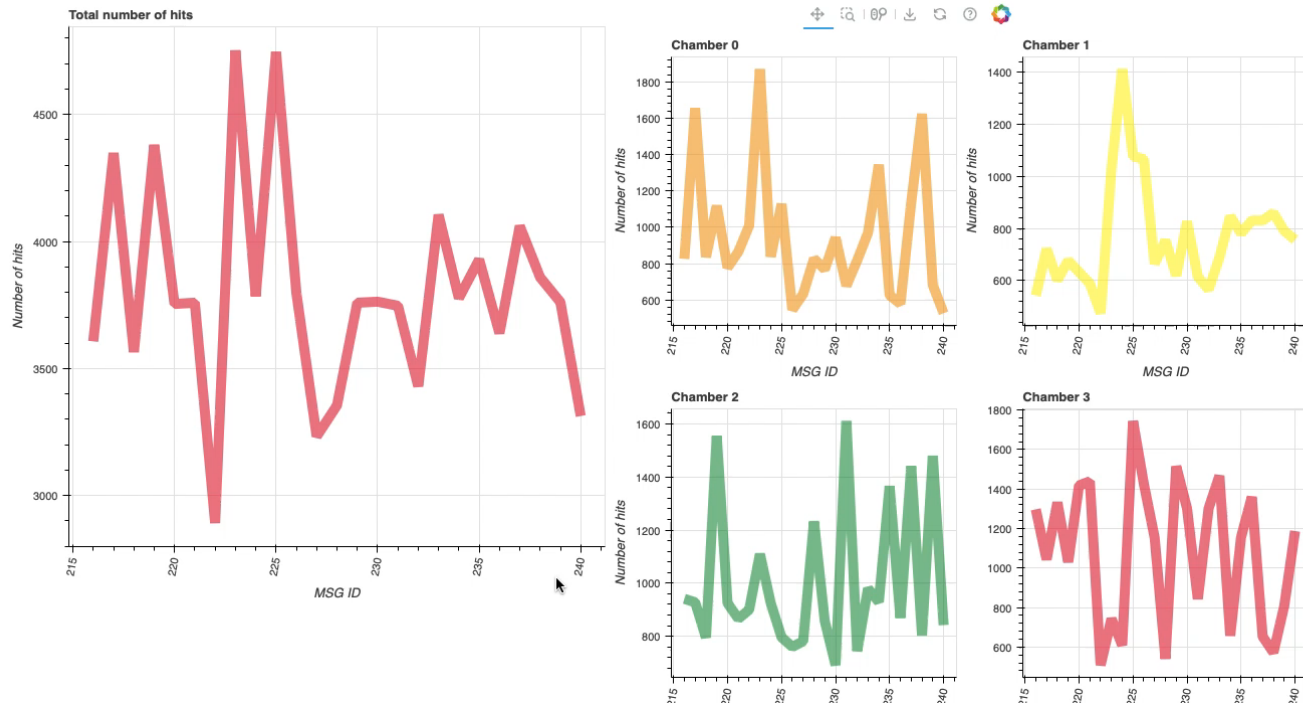
```
curdoc().add_root(grid_with_title)  
curdoc().add_periodic_callback(update,500)
```

- Launch the dashboard → `bokeh serve --show consumer_dashboard.ipynb`

Cosmic Rays Analysis

Diego Bonato, Martina Cacciola, Gabriele Bertinelli, RoBen Bhatti

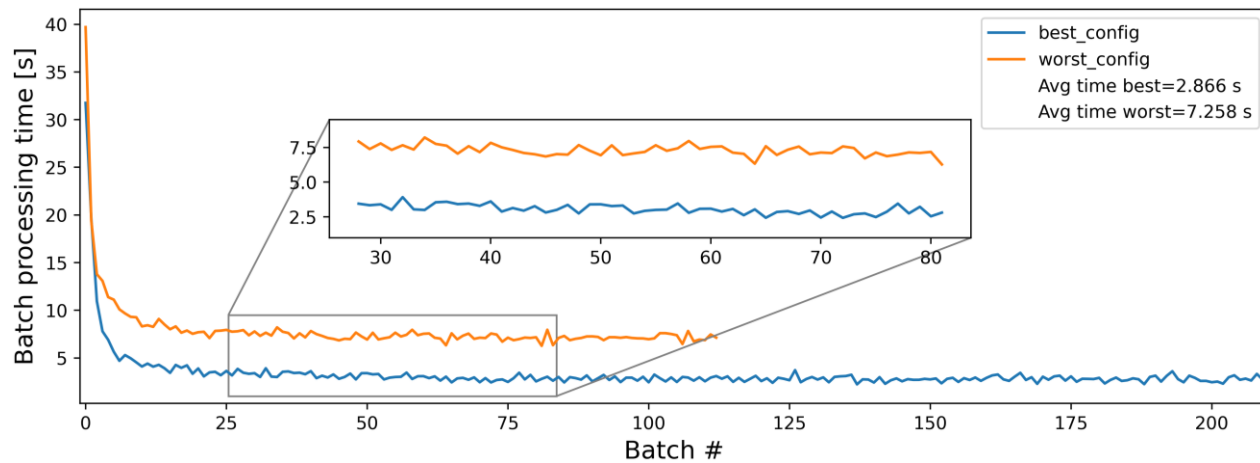

Total number of processed hits

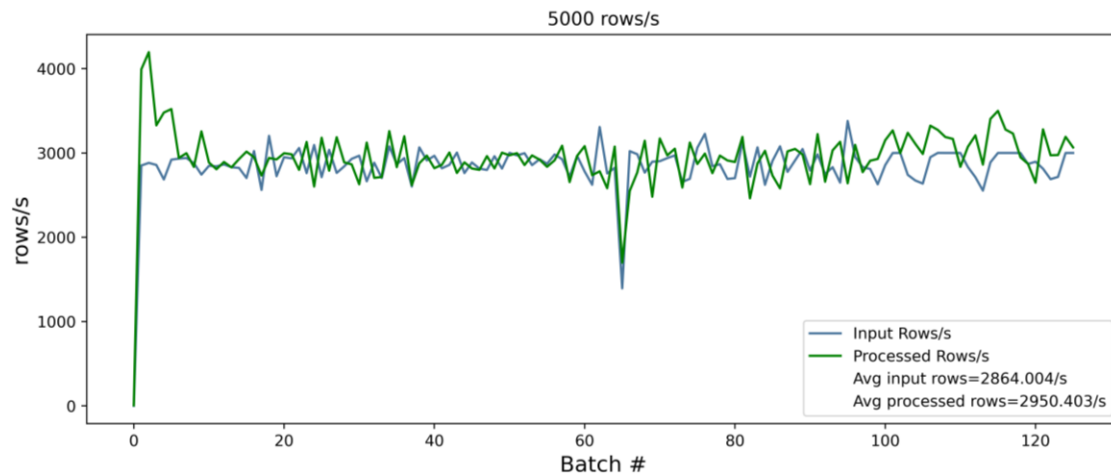
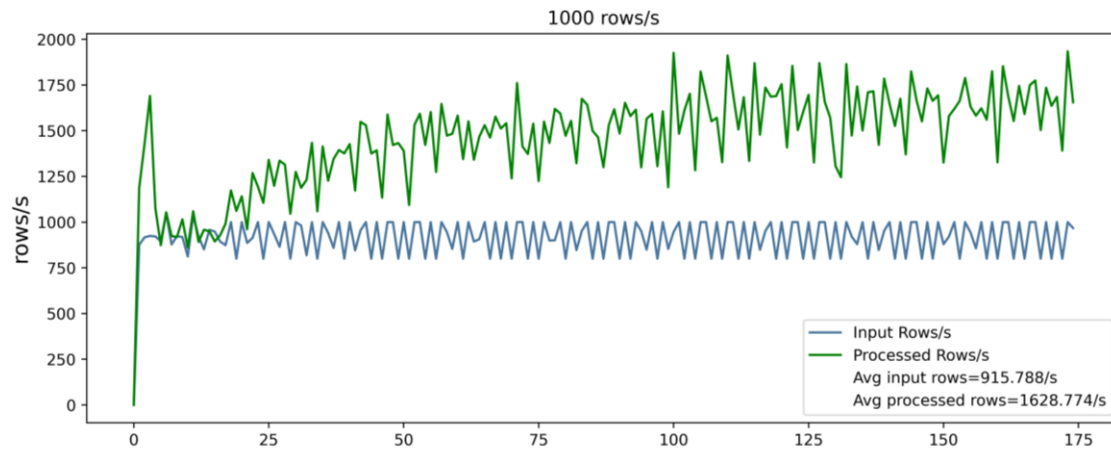


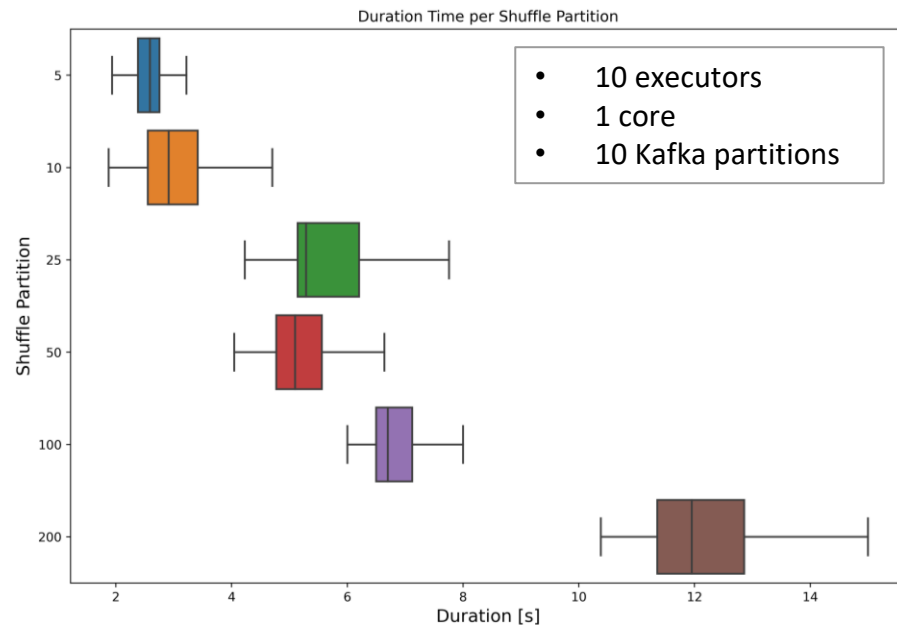
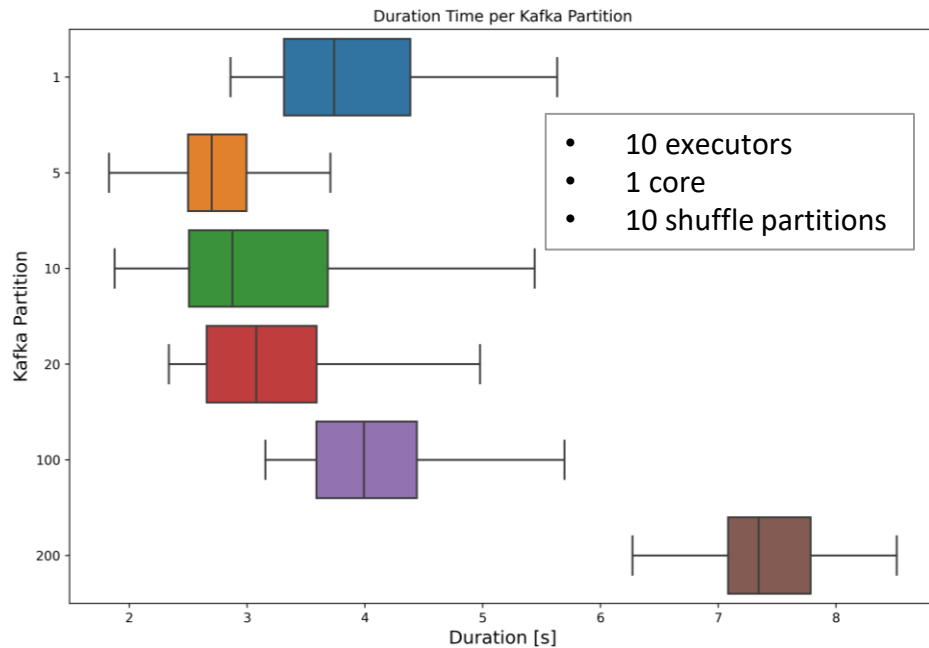
Metrics analysis

We conducted a study to examine how the performance of the network scales by altering the following parameters:


- Number of executors (and cores) → **10 (1 core)**, 3 (2, 4, 4 cores), 5 (2 cores)
- Number of Kafka partitions → 1, 5, **10**, 100, 200
- Number of shuffle partitions → 5, **10**, 25, 50, 100, 200
- Input rows → **1000**, 5000 rows/s







MY HOBBY:



RUNNING A MASSIVE DISTRIBUTED
COMPUTING PROJECT THAT SIMULATES
TRILLIONS AND TRILLIONS OF
TAMAGOTCHIS AND KEEPS THEM
ALL CONSTANTLY FED AND HAPPY

Credit: XKCD



UNIVERSITY OF PADOVA

Department of Physics and Astronomy 'Galileo Galilei'

MSc in Physics of Data